

1. The project goal was to use machine learning to identify POI's (Persons of interest) in the Enron data set. The data set included 146 data points with 21 features for each individual. Some of the financial features were salary, bonuses, and stock options. There were also email features that included number of emails sent to and received from POIs. There were 18 POIs identified in the dataset. Numerous employees were hard to make decisions about due to having "NaN" values.

Several employees had extremely high bonuses and salaries along with negative stock values and holdings. These employees were the outliers. While exploring the data set there were two obviously phony data points - "Total" was included for all employees, and one entry labeled "THE TRAVEL AGENCY IN THE PARK" which proved to not be an employee. Both entries were removed. When creating the new features, I looked at the minimum and maximum values before and after.

```
data_dict.pop('THE TRAVEL AGENCY IN THE PARK')
```

```
data_dict.pop('TOTAL')
```

2. The features I finally settled on are the following:

```
['poi',  
 'exercised_stock_options',  
 'total_stock_value',  
 'bonus',  
 'salary',  
 'percent_to_poi',  
 'deferred_income',  
 'long_term_incentive',  
 'restricted_stock',  
 'total_payments',  
 'shared_receipt_with_poi']
```

I selected these features with SelectKBest which was imported from sklearn.feature_selection module. Feature scaling is an important preprocessing technique to help normalize the dataset and yield balanced features. Imbalanced features cause situations where one feature dominates and overshadows the other, which can compromise the trustworthiness of the experiment. The features that I created were 'from_poi_to_this_person_ratio', and 'from_this_person_to_poi_ratio'. The thinking behind this decision was if the ratio of messages to and from a POI was high, then maybe that person is also a POI. I tested both of these features along with the DecisionTree algorithm, the results were accuracy .82, precision .35, and recall .73.

The feature scores obtained from using SelectKbest are shown below.

```

'SelectKBest scores: '
[('exercised_stock_options', 24.815079733218194),
 ('total_stock_value', 24.18289867856688),
 ('bonus', 20.792252047181535),
 ('salary', 18.289684043404513),
 ('percent_to_poi', 16.40971254803579),
 ('deferred_income', 11.458476579280369),
 ('long_term_incentive', 9.922186013189823),
 ('restricted_stock', 9.2128106219771),
 ('total_payments', 8.772777730091676),
 ('shared_receipt_with_poi', 8.589420731682381),
 ('loan_advances', 7.184055658288725),
 ('expenses', 6.094173310638945),
 ('from_poi_to_this_person', 5.243449713374958),
 ('other', 4.187477506995375),
 ('percent_from_poi', 3.128091748156719),
 ('from_this_person_to_poi', 2.382612108227674),
 ('director_fees', 2.1263278020077054),
 ('to_messages', 1.6463411294420076),
 ('deferral_payments', 0.2246112747360099),

 ('from_messages', 0.16970094762175533),
 ('restricted_stock_deferred', 0.06549965290994214)]

```

In my code I use the default of $k = 5$ and `f_classif` as the score function. The value $k = 5$ is an educated guess, while `chi2`, `f_classif` and `mutual_info_classif` yield a constant score for all the features, `f_classif` produces the reasonable scores shown above.

3. Using an idea from one of my references I created functions that could be easily called and modified to test the data. The algorithm I chose to use was Decision Tree. I also tried using Adaboost, SVM, and Naïve Bayes. SVM took too long to run. The other algorithms were either above or below the .3 threshold for precision and the accuracy was closer to 80%. The final output for all algorithms is:

Adaboost

training time: 0.054 s

predicting time: 0.004 s

accuracy = 0.8604651162790697

Decision Tree

training time: 0.003 s

predicting time: 0.001 s

accuracy = 0.9069767441860465

NaiveBayes

training time: 0.003 s

predicting time: 0.001 s

accuracy = 0.8837209302325582

SVM_CLF

ran a about 1.5 minutes with no output

4. Algorithm tuning is critical to model performance. Limitations can be set on particular behavior of the algorithm, such as the maximum height of a building, or the maximum numbers of estimators. Adaboost terminates boosters. SVM can be used to define kernel types, which allows for the selection - a particular separation of groups that can be more or less linear. **Tuning is the process of seeing the values of parameters for a learning algorithm either by using model defaults, grid search, random search or Bayesian optimization. Tunable parameters of models can affect accuracy, without tuning the parameters well, the model will not optimally solve the machine learning problem. The algorithm that I used was DecisionTree, which is studying the arrays of tunable parameters for decision making. I found no reason to modify the default values, therefore I used the default values without further tuning.**

With Adaboost, the parameters I tuned were:

n_estimators=23, learning_rate=1, algorithm="SAMME"

The approach I used for tuning the parameters was manual intelligent search.

The decision tree algorithm, with the features and modified data set, had the following outcome.

Accuracy: 0.82560

Precision: 0.35333

Recall: 0.37100

F1: 0.36195

F2: 0.36733

Total predictions: 15000

True positives: 742

False positives: 1358

False negatives: 1258

True negatives: 11642

5. Validation in Machine Learning is the process used to check the validity of the training model by testing the model on an unseen dataset. A classic error when validating data would be to train your model on all of the data at one time. To validate my data, I used the `stratifiedshufflesplit` found in the `tester.py` script provided by Udacity.

6. Using the `stratifiedshufflesplit` my final model had these average performance metrics:

Accuracy: 0.83053

This accuracy means the model was 83.1% accurate in the prediction of whether or not a person was a POI.

Precision: 0.36382

This means that 36.4% of people the model classified as POIs were really POIs.

Recall: 0.36200

This means that the model correctly identified 36.2% of the POIs in the data.